

DS2_Assignment1_NF

March 20, 2024

1 DS2 Assignment 1 - Nicolas Fernandez

1.1 Predicting Property Prices from the Xindian District of New Taipei City, Taiwan

The task is to predict property prices using data taken from [UC Irvine](#) in order to build a web app where buyers/sellers could rate their homes.

The precise dataset being used is a cleaned version uploaded within Janos Divenyi's [github repository](#).

The task asks to create a 20% subsample from the data and then create a 70/30% split from that subsample for a training and test sets.

1.2 Reading Data and Creating Training and Test Splits

```
[1]: # Importing required libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

# Setting the pseudo random seed for the due date of the assignment
prng = np.random.RandomState(20240322)

# Reading the data
#real_estate_data = pd.read_csv("https://raw.githubusercontent.com/divenyi/janos/
↪ceu-ml/2023/data/real_estate/real_estate.csv")
real_estate_data = pd.read_csv('real_estate.csv')
display(real_estate_data.head())
print(real_estate_data.info())
```

	id	transaction_date	house_age	distance_to_the_nearest_MRT_station	\
0	1	2012.917	32.0	84.87882	
1	2	2012.917	19.5	306.59470	
2	3	2013.583	13.3	561.98450	
3	4	2013.500	13.3	561.98450	
4	5	2012.833	5.0	390.56840	

	number_of_convenience_stores	latitude	longitude	house_price_of_unit_area
0	10	24.98298	121.54024	37.9

1	9	24.98034	121.53951	42.2
2	5	24.98746	121.54391	47.3
3	5	24.98746	121.54391	54.8
4	5	24.97937	121.54245	43.1

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 414 entries, 0 to 413
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	id	414 non-null	int64
1	transaction_date	414 non-null	float64
2	house_age	414 non-null	float64
3	distance_to_the_nearest_MRT_station	414 non-null	float64
4	number_of_convenience_stores	414 non-null	int64
5	latitude	414 non-null	float64
6	longitude	414 non-null	float64
7	house_price_of_unit_area	414 non-null	float64

```
dtypes: float64(6), int64(2)
```

```
memory usage: 26.0 KB
```

```
None
```

From the data there's a superfluous `id` column that will be dropped. The `transaction_date` column contains information about the year and month of the observation however it is in a non-standard format (`pd.to_datetime()` cannot be used) and will not be included in initial features. The `latitude` and `longitude` columns contains geographical data for each column but that cannot be made sense of without further information and therefore will not be included for initial feature modeling.

```
[2]: # Dropping the id column because it's unnecessary
real_estate_data.drop(columns=['id'], inplace=True)

# Creating a randomly selected 20% sample of the data, creating target variable
↳ 'outcome'
real_estate_sample = real_estate_data.sample(frac=0.2, random_state=prng)
outcome = real_estate_sample["house_price_of_unit_area"]

# Selecting numeric features as X variables, excluding target variable,
↳ transaction_date, latitude, and longitude
features = real_estate_sample.drop(columns=["house_price_of_unit_area",
↳ "transaction_date", "latitude", "longitude"])

# Creating training and testing splits from features and target, using a 30%
↳ split per specifications
X_train, X_test, y_train, y_test = train_test_split(features, outcome,
↳ test_size=0.3, random_state=prng)
```

```
print(f"Size of the training set: {X_train.shape}, size of the test set: {X_test.shape}")
X_train.head()
```

Size of the training set: (58, 3), size of the test set: (25, 3)

```
[2]:      house_age  distance_to_the_nearest_MRT_station  \
331         25.6                                4519.69000
386          0.0                                185.42960
267         34.7                                482.75810
252          5.9                                90.45606
197         34.4                                126.72860

      number_of_convenience_stores
331                               0
386                               0
267                               5
252                               9
197                               8
```

1.3 Evalaution Function

```
[3]: # Defining the loss function, using Root Mean Squared Log Error
def calculateRMSLE(prediction, y_obs):
    return round(np.sqrt(np.mean((np.log(np.where(prediction < 0, 0, prediction) + 1) - np.log(y_obs + 1))**2)), 4)
```

RMSLE is an appropriate loss function since its a calculation that is less sensitive to outliers by design, appropriate for property price prediction. The risk (from a business perspective) from making a wrong prediction could be either under or over pricing homes in certain areas because RMSLE may treat values as outliers that may have be more important than they seem. Overall, however, property prices are susceptible to large increases based on several factors and can increase in price in orders of magnitude rather than a more gradual scale and for this reason using a loss function on a logarithmic scale is preferred.

1.4 Rigid Models

Initially OLS models without any non-linearity will be created

1.4.1 Benchmark Model

```
[4]: # Creating a simple benchmark model for comparison, mean value of target variable for training set
benchmark = np.mean(y_train)

# Calculating performance of benchmark on test set using loss function
benchmark_pred = ['Benchmark', calculateRMSLE(benchmark, y_train), calculateRMSLE(benchmark, y_test)]
```

```
# Storing and displaying results in a dataframe
results = pd.DataFrame([benchmark_pred], columns=['Model', 'Train', 'Test'])
results
```

```
[4]:      Model  Train  Test
0  Benchmark  0.3434  0.3221
```

1.4.2 OLS Single Feature - distance_to_the_nearest_MRT_station

```
[5]: # Importing required library
from sklearn.linear_model import LinearRegression

# Creating OLS model trained on only `distance_to_the_nearest_MRT_station`
ols_single = LinearRegression().
    ↪fit(X_train[['distance_to_the_nearest_MRT_station']], y_train)

# Creating predictions for model and calculating RMSLE
train_error = calculateRMSLE(ols_single.
    ↪predict(X_train[['distance_to_the_nearest_MRT_station']]), y_train)
test_error = calculateRMSLE(ols_single.
    ↪predict(X_test[['distance_to_the_nearest_MRT_station']]), y_test)
ols_single_pred = ['Single Feature OLS', train_error, test_error]

# Adding to results
results.loc[len(results)] = ols_single_pred
results
```

```
[5]:      Model  Train  Test
0      Benchmark  0.3434  0.3221
1 Single Feature OLS  0.2250  0.2305
```

This model improves upon using the average of the target variable as a predictive model for our target as the RMSLE's improves using a single feature from the dataset. The model can be significantly improved as it likely does not catch much of the complexity of the data but using `distance_to_the_nearest_MRT_station`, determined to be the most significant variable to predicting property prices, as an explanatory variable has improved our predictions.

1.4.3 Multivariate OLS

Building an OLS model that uses all the available meaningful features instead of a single variable.

```
[6]: # Creating multivariate OLS model and fitting to training data
ols_multi = LinearRegression().fit(X_train, y_train)

# Creating predictions
train_error = calculateRMSLE(ols_multi.predict(X_train), y_train)
test_error = calculateRMSLE(ols_multi.predict(X_test), y_test)
```

```
ols_multi_pred = ['Multi OLS', train_error, test_error]

# Adding to results
results.loc[len(results)] = ols_multi_pred
results
```

```
[6]:
```

	Model	Train	Test
0	Benchmark	0.3434	0.3221
1	Single Feature OLS	0.2250	0.2305
2	Multi OLS	0.1993	0.2317

The multivariate model using all the available meaningful features improved the predictive power on the training set but performed slightly worse on the test set. Using all the available features does not increase the predictive power over using only a single feature. This implies that the other features in the data (house_age and number_of_convenience_stores) causes a slight overfit compared to only using distance_to_the_nearest_MRT_station and that we may have introduced more noise, bias, and/or variance into our model.

1.5 Flexible Models

Below more flexible models will be created

1.5.1 Polynomial OLS with Interactions

```
[7]: # Importing required library
from sklearn.preprocessing import PolynomialFeatures

# Creating Polynomial object to create polynomial terms and interactions in
↳ training data
poly_interactions = PolynomialFeatures(degree=2, include_bias=False,
↳ interaction_only=True)
X_poly = poly_interactions.fit_transform(X_train)

# Creating polynomial OLS model
ols_poly = LinearRegression().fit(X_poly, y_train)

# Creating predictions
train_error = calculateRMSLE(ols_poly.predict(X_poly), y_train)
test_error = calculateRMSLE(ols_poly.predict(poly_interactions.
↳ transform(X_test)), y_test)
ols_poly_pred = ['Poly OLS', train_error, test_error]

# Adding to results
results.loc[len(results)] = ols_poly_pred
results
```

```
[7]:
```

	Model	Train	Test
0	Benchmark	0.3434	0.3221

1	Single Feature OLS	0.2250	0.2305
2	Multi OLS	0.1993	0.2317
3	Poly OLS	0.1742	0.1611

Adding polynomial and interaction terms to the model has improved the RMSLE on the test set somewhat significantly. It appears that adding these terms to the OLS model and making it more flexible has captured non-linearity in the data. Also the interaction terms may be reducing bias in the model which would also increase the performance on the test set.

1.5.2 Feature Engineering

In order to improve the model further feature engineering will need to be done. This means making sense of the transaction dates and geographical data.

From research, New Taipei City is actually within Taipei City and used to be known as Taipei County but was renamed because it exceeded a population of 2 million which meant that it had to be redefined as it's own city by law. Xindian is a county within the redefined Taipei County. For the purposes of this analysis the city center for Taipei City (not New Taipei) will be considered the city center for which all observations will be compared to. The geographical data will then be converted to km distance from the city center based on the provided coordinates and added to the dataframe as a new column. The latitude for the city center that will be used is 25.03583333 and the longitude is 121.5683333. The information was pulled from [here](#) and verified with Google Maps.

The transaction dates within the data follow the format 2013.500, for example, with the number to the left of the decimal being the year and the number to the right of the decimal being the decimal representation of the month (.500 = June = 6/12). The only quirk to this is that .000 = December. This data will be converted and split into two columns, one for the year and another for the month.

These features are likely to be important because from general domain knowledge property prices tend to increase the closer they are to the geographical city center of the respective major city and also property prices from year to year tend to increase as well, sometimes due more to general trends in the market and also sometimes due to general inflation. Enabling this data to be used will likely prove beneficial for model building.

```
[8]: # Defining function for creating year and month columns from transaction_date
      ↪as categorical columns
def date_fix(df):
    df['year'] = df['transaction_date'].astype(int)
    df['year'] = df['year'].astype('category')

    df['month'] = np.where(df['transaction_date'] % 1 == 0, 12,
                           round((df['transaction_date'] % 1) * 12).astype(int))
    df['month'] = df['month'].astype('category')

# Defining function for calculating distance from city center in km
def cc_distance(lat, lon):
    # City center coordinates (Taipei)
    cc_lat = 25.03583333
```

```

cc_lon = 121.5683333

# Convert latitude and longitude from degrees to radians
lat, lon, cc_lat, cc_lon = np.radians([lat, lon, cc_lat, cc_lon])

# Calculating difference from city center
diff_lon = lon - cc_lon
diff_lat = lat - cc_lat

# Using the Haversine Formula to calculate distance from radians
a = np.sin(diff_lat / 2) ** 2 + np.cos(lat) * np.cos(cc_lat) * np.
↪ sin(diff_lon / 2) ** 2
c = 2 * np.arctan2(np.sqrt(a), np.sqrt(1 - a))
km_distance = 6371 * c # 6371 = Radius of Earth in km

return km_distance

# Running feature engineering functions on sample
date_fix(real_estate_sample)
real_estate_sample['km_distance_from_cc'] = real_estate_sample.apply(lambda row:
↪ cc_distance(row['latitude'], row['longitude']), axis=1)

# Creating feature engineered matrix and saving column names to a list
feature_matrix = real_estate_sample.drop(columns=['house_price_of_unit_area',
↪ 'transaction_date', 'latitude', 'longitude'])

# Creating new data split with feature engineered data with same specifications
X_train_fe, X_test_fe, y_train_fe, y_test_fe = train_test_split(feature_matrix,
↪ outcome, test_size=0.3, random_state=prng)

print(f"Size of the FE training set: {X_train_fe.shape}, size of the FE test_
↪ set: {X_test_fe.shape}")
X_train_fe.head()

```

Size of the FE training set: (58, 6), size of the FE test set: (25, 6)

```

[8]:      house_age  distance_to_the_nearest_MRT_station  \
209         34.8                                175.62940
361         41.4                                281.20500
81          30.8                                377.79560
114         30.6                                143.83830
160          3.5                                 49.66105

      number_of_convenience_stores  year  month  km_distance_from_cc
209                             8  2012     10           7.399630
361                             8  2013      1           7.466184
81                              6  2013     12           8.466577

```

114	8	2012	8	6.617315
160	8	2012	11	9.155881

1.5.3 OLS Model with Feature Engineered Data

```
[9]: # Importing required libraries
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Setting up OneHotEncoder for handling categorical variables 'year' and 'month'
one_hot_encoder = OneHotEncoder(sparse_output=False, drop="first")
categorical_vars = ['year', 'month']

# Using ColumnTransformer to create the dummies based on defined OneHotEncoder
column_transformer = ColumnTransformer([("create_dummies", one_hot_encoder,
    categorical_vars)],

    remainder="passthrough")

# Creating OLS model for FE data with Pipeline
pipe_ols_fe = Pipeline([("preprocess", column_transformer),
    ("ols", LinearRegression())])

# Fitting the data
pipe_ols_fe.fit(X_train_fe, y_train_fe)

# Creating predictions
train_error = calculateRMSLE(pipe_ols_fe.predict(X_train_fe), y_train_fe)
test_error = calculateRMSLE(pipe_ols_fe.predict(X_test_fe), y_test_fe)
ols_fe_pred = ['FE OLS', train_error, test_error]

# Adding to results
results.loc[len(results)] = ols_fe_pred
results
```

```
[9]:
```

	Model	Train	Test
0	Benchmark	0.3434	0.3221
1	Single Feature OLS	0.2250	0.2305
2	Multi OLS	0.1993	0.2317
3	Poly OLS	0.1742	0.1611
4	FE OLS	0.1733	0.2043

The OLS model with feature engineering performed worse on the test set than the non-feature engineered Poly OLS model but is still performing better than the more rigid OLS models without the feature engineered data. This model is not accounting for any non-linearity in the data and also not adding interaction terms and is not producing as good of a fit on the test set as a result compared to the polynomial OLS on the base features.

1.5.4 Polynomial OLS with Feature Engineered Data

```
[10]: # Importing required library
from sklearn.feature_selection import VarianceThreshold # To account for
    ↪ categorical values with zero variance

# Initiating VarianceThreshold to drop categorical variables with no variance
drop_no_variance = VarianceThreshold()

# Creating model with Pipeline
pipe_ols_poly_fe = Pipeline([
    ('preprocess', column_transformer),
    ('interactions', poly_interactions),
    ('drop_zero_variance', drop_no_variance),
    ('ols', LinearRegression())
])

# Fitting the data
pipe_ols_poly_fe.fit(X_train_fe, y_train_fe)

# Creating predictions
train_error = calculateRMSLE(pipe_ols_poly_fe.predict(X_train_fe), y_train_fe)
test_error = calculateRMSLE(pipe_ols_poly_fe.predict(X_test_fe), y_test_fe)
ols_poly_fe_pred = ['FE Poly OLS', train_error, test_error]

# Adding to results
results.loc[len(results)] = ols_poly_fe_pred
results
```

```
[10]:
```

	Model	Train	Test
0	Benchmark	0.3434	0.3221
1	Single Feature OLS	0.2250	0.2305
2	Multi OLS	0.1993	0.2317
3	Poly OLS	0.1742	0.1611
4	FE OLS	0.1733	0.2043
5	FE Poly OLS	0.0479	0.8053

The polynomial and interaction terms included in the OLS model using the feature engineered data performed better on the training set but performed terribly on the test set. There is a clear overfit when adding interactions and polynomial terms on the feature engineered data to account for non-linearity. This implies that while there is some non-linearity in the data, it clearly relates to the features from the base features rather than the feature engineered data.

1.5.5 Random Forest

```
[11]: # Importing required library
from sklearn.ensemble import RandomForestRegressor

# Creating a Random Forest model with Pipeline
pipe_rf = Pipeline([('preprocess', column_transformer),
                    ("rf", RandomForestRegressor(random_state=prng))])

# Fitting the model to the feature engineered data
pipe_rf.fit(X_train_fe, y_train_fe)

# Creating predictions
train_error = calculateRMSLE(pipe_rf.predict(X_train_fe), y_train_fe)
test_error = calculateRMSLE(pipe_rf.predict(X_test_fe), y_test_fe)
rf_pred = ['RF', train_error, test_error]

# Adding to results
results.loc[len(results)] = rf_pred
results
```

```
[11]:
```

	Model	Train	Test
0	Benchmark	0.3434	0.3221
1	Single Feature OLS	0.2250	0.2305
2	Multi OLS	0.1993	0.2317
3	Poly OLS	0.1742	0.1611
4	FE OLS	0.1733	0.2043
5	FE Poly OLS	0.0479	0.8053
6	RF	0.0844	0.1545

The Random Forest model does the best job so far of creating a predictive model. This can be attributed to the nature of a Random Forest model using bagging and regularization to become more robust to noise within the data since it averages predictions among many trees, in this case 100 (the default value). The feature importance decided by the algorithm is also doing a better job of capturing the relationships in the data to the target variable.

Below a table is shown with the feature importance of the Random Forest model created from the feature engineered data.

```
[12]: # Creating dataframe for readability for displaying variable importance
rf_var_imp = pd.DataFrame(
    pipe_rf['rf'].feature_importances_,
    pipe_rf[:-1].get_feature_names_out())\
    .reset_index()\
    .rename({"index": "variable", 0: "imp"}, axis=1)\
    .sort_values(by=["imp"], ascending=False)\
    .reset_index(drop = True)

# Creating cumulative sum column
```

```

rf_var_imp['cumulative_imp'] = rf_var_imp['imp'].cumsum()

# Displaying dataframe with formatting
rf_var_imp.style.format({'imp': lambda x: f'{x:,.1%}', 'cumulative_imp': lambda x: f'{x:,.1%}'})

```

[12]: <pandas.io.formats.style.Styler at 0x12537c4bc50>

From these results we can see that `distance_to_the_nearest_MRT_station` is in fact the most important feature in the data, as theorized earlier. The next three significant features (although significantly less so than MRT station distance) are `km_distance_from_cc`, `house_age`, and `number_of_convenience_stores`. This makes sense as people in general value properties more that are closer to the city center and the older the house usually the less expensive it tends to be. It also appears that a house being closer to more convenience stores also influences the price of properties in Taipei which implies that people in that cit value having better access to convenience stores.

Conversely, the date of the purchase appears to have little to do with making a good prediction. This could be that since there isn't much variance in the amount of time within the data that the transaction date has little variance.

Below we take a look at a single tree to illustrate:

```

[13]: # Importing required libraries
from sklearn import tree
import matplotlib.pyplot as plt

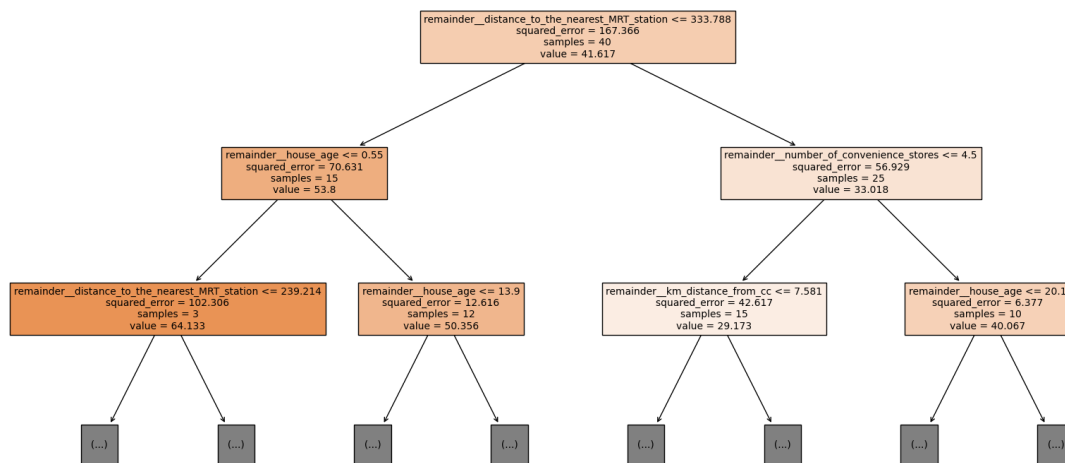
# Look at single trees
chosen_tree = pipe_rf["rf"].estimators_[0]

plt.figure(figsize=(20,10))

# Plot tree with adjusted settings
tree.plot_tree(chosen_tree,
                feature_names=pipe_rf[:-1].get_feature_names_out(),
                max_depth=2, # Plotting only a depth of 2 splits
                fontsize=10, # Adjust font size of text in nodes
                filled=True) # Fill nodes with colors

plt.show()

```



1.5.6 XGBoost Model

```
[14]: # Importing required library
import xgboost as xgb

# Creating XGB model, fitting to feature engineered training data
xgb_fe = xgb.XGBRegressor(enable_categorical=True).fit(X_train_fe, y_train_fe)

# Creating predictions
train_error = calculateRMSLE(xgb_fe.predict(X_train_fe), y_train_fe)
test_error = calculateRMSLE(xgb_fe.predict(X_test_fe), y_test_fe)
xgb_pred = ['XGB', train_error, test_error]

# Adding to results
results.loc[len(results)] = xgb_pred
results
```

```
[14]:
```

	Model	Train	Test
0	Benchmark	0.3434	0.3221
1	Single Feature OLS	0.2250	0.2305
2	Multi OLS	0.1993	0.2317
3	Poly OLS	0.1742	0.1611
4	FE OLS	0.1733	0.2043
5	FE Poly OLS	0.0479	0.8053
6	RF	0.0844	0.1545
7	XGB	0.0000	0.1736

```
[15]: # Viewing features from XGBoost model to see how it split the data, which
      ↪categories it used, etc.
      xgb_fe.get_booster().trees_to_dataframe().head()
```

```
[15]:
```

	Tree	Node	ID	Feature	Split	Yes	\
0	0	0	0-0	distance_to_the_nearest_MRT_station	665.063599	0-1	
1	0	1	0-1	distance_to_the_nearest_MRT_station	337.601593	0-3	
2	0	2	0-2	km_distance_from_cc	12.772239	0-5	
3	0	3	0-3	month	NaN	0-8	
4	0	4	0-4	house_age	3.500000	0-9	

	No Missing		Gain	Cover	Category
0	0-2	0-2	4520.014650	58.0	None
1	0-4	0-4	937.893066	39.0	None
2	0-6	0-6	285.436523	19.0	None
3	0-7	0-7	308.360840	24.0	[2, 3, 4, 5, 8, 11]
4	0-10	0-10	141.103470	15.0	None

From the results of the XGB we can see that the model performs perfectly on the training set but performs worse on the test set than the RF model. The XGBoost model is still doing a better job of predicting on the test set than most models but is also worse than the Poly OLS model (base features). The interpretation of this is that the XGBoost model is capturing some (and maybe most) but not all of the non-linearity in the data and is therefore performing worse than the two models mentioned on the test set. It could also be that some of the features used may be adding more noise/bias/variance to the data rather than helping and XGBoost may not be doing as good of a job of accounting for this as the Random Forest model.

1.6 Initial Analysis

From the models constructed so far, it would appear that the Random Forest model is our best choice for going forward and creating a web app designed to aid buyers/sellers to rate their homes. As it stands, however, I would like to test adding more data to the training set to see how the models fare and if they improve or worsen. Along with this, a potential more careful, theoretical approach of choosing polynomial and interaction terms could yield a better OLS model than the Poly OLS that is currently constructed. Some more options for potential improvement are cross-validation, hyperparameter tuning (e.g. adding a grid search to the RF), or using another model type altogether.

1.7 Rerunning Models on Full Dataset, Excluding Test Sets

Below several of the models created (both rigid and flexible) will be tested on the overall dataset. We will be careful to give the models trained on the full dataset the same comparison base and not include that test data in the training set for both the base and feature engineered variables.

```
[16]: # Creating the training set for the base features models
      real_estate_full = real_estate_data.loc[~real_estate_data.index.isin(X_test.
      ↪index)]
      y_train_full = real_estate_full['house_price_of_unit_area']
```

```

X_train_full = real_estate_full.drop(columns=["house_price_of_unit_area",
↳ "transaction_date", "latitude", "longitude"])

# Doing the same for the feature engineered set
date_fix(real_estate_data)
real_estate_full_fe = real_estate_data.loc[~real_estate_data.index.
↳ isin(X_test_fe.index)].copy()
real_estate_full_fe['km_distance_from_cc'] = real_estate_full_fe.apply(lambda
↳ row: cc_distance(row['latitude'], row['longitude']), axis=1)
y_train_full_fe = real_estate_full_fe['house_price_of_unit_area']
X_train_full_fe = real_estate_full_fe.drop(columns=['house_price_of_unit_area',
↳ 'transaction_date', 'latitude', 'longitude'])

# Checking sample size results
print(f'Size of original dataset: {real_estate_data.shape[0]}')
print(f'Shape of base features full training set: {X_train_full.shape}, shape
↳ of base features test set: {X_test.shape}')
print(f'Shape of feature engineered full training set: {X_train_full_fe.shape},
↳ shape of feature engineered test set: {X_test_fe.shape}')

```

Size of original dataset: 414

Shape of base features full training set: (389, 3), shape of base features test set: (25, 3)

Shape of feature engineered full training set: (389, 6), shape of feature engineered test set: (25, 6)

```

[17]: # Running the single OLS on the full training set
ols_single_full = LinearRegression().
↳ fit(X_train_full[['distance_to_the_nearest_MRT_station']], y_train_full)

# Creating predictions for model and calculating RMSLE
train_error = calculateRMSLE(ols_single_full.
↳ predict(X_train_full[['distance_to_the_nearest_MRT_station']] ), y_train_full)
test_error = calculateRMSLE(ols_single_full.
↳ predict(X_test[['distance_to_the_nearest_MRT_station']] ), y_test)
ols_single_pred_full = ['Single Feature OLS Full', train_error, test_error]

# Adding to results
results.loc[len(results)] = ols_single_pred_full

# Running the base features Polynomial OLS on the full training set
X_poly_full = poly_interactions.fit_transform(X_train_full)

# Creating polynomial OLS model
ols_poly_full = LinearRegression().fit(X_poly_full, y_train_full)

# Creating predictions

```

```

train_error = calculateRMSLE(ols_poly_full.predict(X_poly_full), y_train_full)
test_error = calculateRMSLE(ols_poly_full.predict(poly_interactions.
    ↪transform(X_test)), y_test)
ols_poly_pred_full = ['Poly OLS Full', train_error, test_error]

# Adding to results
results.loc[len(results)] = ols_poly_pred_full

# Running the Random Forest model on the full training set
pipe_rf_full = Pipeline([('preprocess', column_transformer),
    ("rf", RandomForestRegressor(random_state=prng))])

# Fitting the model to the feature engineered data
pipe_rf_full.fit(X_train_full_fe, y_train_full_fe)

# Creating predictions
train_error = calculateRMSLE(pipe_rf_full.predict(X_train_full_fe),
    ↪y_train_full_fe)
test_error = calculateRMSLE(pipe_rf_full.predict(X_test_fe), y_test_fe)
rf_pred_full = ['RF Full', train_error, test_error]

# Adding to results
results.loc[len(results)] = rf_pred_full
results

```

```

[17]:
      Model  Train  Test
0      Benchmark  0.3434  0.3221
1  Single Feature OLS  0.2250  0.2305
2      Multi OLS  0.1993  0.2317
3      Poly OLS  0.1742  0.1611
4      FE OLS  0.1733  0.2043
5  FE Poly OLS  0.0479  0.8053
6      RF  0.0844  0.1545
7      XGB  0.0000  0.1736
8  Single Feature OLS Full  0.3477  0.2211
9      Poly OLS Full  0.2531  0.1677
10     RF Full  0.0747  0.1682

```

The models chosen to be trained on the full set were the least flexible single feature OLS, the flexible polynomial OLS using base features, and the Random Forest. From the results we can see that the single feature OLS did improve on the test set but that both the Poly OLS and RF models performed worse on the test set when trained on the full training set instead of a sample training set. Both of these models still perform better than the single feature OLS and between them the RF model worsened to the point where it performs very slightly worse than the Poly OLS model now.?

1.8 Final Analysis

While typically more data tends to improve and regularize any potentially extraneous data points in smaller sample sizes it would appear that using more data to train the models is improving the predictions. Having seen how the best models from sample training set performed worse (though not significantly so) it gives me less confidence that I should launch the web app. That being said I would still feel confident that launching the web app using the RF model would be the way forward as the RMSLE values on the test set are still very low and further tweaking of that model could yield even better results.